

# ResNet-18 PyTorch Lightning



I want to build a ResNet-18 for image classification. My input data are images, but they are different size. I need to resize them to 256\*256. The output should be a probability of one of ten classes. Use PyTorch to code it and use PyTorch Lightning to build trainer, upload learning metrics and loss to Weights & Biases.



Here's a step-by-step guide on how to build a ResNet-18 model for image classification with PyTorch and PyTorch Lightning, including resizing the input images to 256x256 and logging metrics and losses to Weights & Biases.

## Step 1: Install the necessary libraries

Make sure you have the following libraries installed:

```
bash
```

```
pip install torch torchvision pytorch-lightning wandb
```

## Step 2: Import the necessary libraries

```
python
```

```
import torch
import torch.nn as nn
import torchvision.transforms as transforms
import torchvision.models as models
import pytorch_lightning as pl
from pytorch_lightning.loggers import WandbLogger
from torch.utils.data import DataLoader, random_split
from torchvision.datasets import ImageFolder
```

## Step 3: Define the Data Module

We'll create a PyTorch Lightning Data Module to handle data loading and transformations.

```
python
```

```
class ImageDataModule(pl.LightningDataModule):
    def __init__(self, data_dir, batch_size=32, image_size=256):
        super().__init__()
        self.data_dir = data_dir
        self.batch_size = batch_size
        self.image_size = image_size
        self.transform = transforms.Compose([
            transforms.Resize((self.image_size, self.image_size)),
```

```

        transforms.ToTensor(),
        transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225))
    ])

    def setup(self, stage=None):
        dataset = ImageFolder(self.data_dir, transform=self.transform)
        self.train_data, self.val_data = random_split(dataset, [int(len(dataset)*0.8),
len(dataset) - int(len(dataset)*0.8)])

    def train_dataloader(self):
        return DataLoader(self.train_data, batch_size=self.batch_size, shuffle=True)

    def val_dataloader(self):
        return DataLoader(self.val_data, batch_size=self.batch_size)

```

## Step 4: Define the Model

We'll define the model using PyTorch and wrap it with a PyTorch Lightning module.

python

```

class ResNetClassifier(pl.LightningModule):
    def __init__(self, num_classes=10):
        super(ResNetClassifier, self).__init__()
        self.model = models.resnet18(pretrained=True)
        self.model.fc = nn.Linear(self.model.fc.in_features, num_classes)
        self.criterion = nn.CrossEntropyLoss()

    def forward(self, x):
        return self.model(x)

    def training_step(self, batch, batch_idx):
        images, labels = batch
        outputs = self(images)
        loss = self.criterion(outputs, labels)
        self.log('train_loss', loss)
        return loss

    def validation_step(self, batch, batch_idx):
        images, labels = batch
        outputs = self(images)
        loss = self.criterion(outputs, labels)
        acc = (outputs.argmax(dim=1) == labels).float().mean()
        self.log('val_loss', loss, prog_bar=True)
        self.log('val_acc', acc, prog_bar=True)

    def configure_optimizers(self):
        return torch.optim.Adam(self.parameters(), lr=1e-4)

```

## Step 5: Train the Model

We'll use PyTorch Lightning's trainer to train the model and log the metrics to Weights & Biases.

python

```

if __name__ == "__main__":
    import os
    import wandb
    from pytorch_lightning.callbacks import ModelCheckpoint

    # Initialize Weights & Biases
    wandb.init(project="image-classification")

    # Define the data module

```

```
data_dir = 'path_to_your_dataset'
data_module = ImageDataModule(data_dir)

# Define the model
model = ResNetClassifier(num_classes=10)

# Define the logger
wandb_logger = WandbLogger(project="image-classification")

# Define the trainer
trainer = pl.Trainer(
    max_epochs=10,
    gpus=1 if torch.cuda.is_available() else 0,
    logger=wandb_logger,
    callbacks=[ModelCheckpoint(monitor='val_loss')]
)

# Train the model
trainer.fit(model, data_module)

# Finish the Weights & Biases run
wandb.finish()
```

## Summary

This code will:

1. Define a PyTorch Lightning Data Module to load and preprocess the images.
2. Define a ResNet-18 model modified for your classification task.
3. Use PyTorch Lightning to handle training and validation.
4. Log training metrics and loss to Weights & Biases.

Make sure to replace `'path_to_your_dataset'` with the actual path to your dataset. This setup assumes you have your dataset organized in subdirectories for each class, as expected by `torchvision.datasets.ImageFolder`.

Feel free to adjust hyperparameters and other settings as needed for your specific use case.



Now, I need to increase the batch size, but my gpu is limited size. I want to use Gradient accumulation. How to use it in pytorch lightningning framework?



In PyTorch Lightning, implementing gradient accumulation is straightforward and can be very useful when you're limited by GPU memory but want to use a larger batch size effectively. Gradient accumulation allows you to accumulate the gradients over several mini-batches and only update the model weights after a specified number of forward and backward passes, effectively increasing the effective batch size without additional memory cost.

## Step 1: Modify the Trainer

To enable gradient accumulation in PyTorch Lightning, you simply need to set the `accumulate_grad_batches` parameter in the `Trainer`. This parameter dictates how many batches to accumulate before performing a backward/update pass.

Here's how you can modify the trainer setup to include gradient accumulation:

python

```
from pytorch_lightning import Trainer

# Define the trainer with gradient accumulation
trainer = pl.Trainer(
    max_epochs=10,
    gpus=1 if torch.cuda.is_available() else 0,
    accumulate_grad_batches=4, # Accumulate gradients over 4 batches
    logger=wandb_logger,
    callbacks=[ModelCheckpoint(monitor='val_loss')]
)
```

## Explanation

In the above code, `accumulate\_grad\_batches=4` means that the optimizer will step every 4 batches, effectively quadrupling the batch size in terms of gradient accumulation without increasing the GPU memory required per batch. This is particularly useful when dealing with very large models or very large datasets.

## Additional Considerations

- **Learning Rate Adjustment:** When using gradient accumulation, the effective batch size becomes larger. Often, it is suggested to increase the learning rate in proportion to the increased effective batch size to maintain the same optimization dynamics.
- **Logging Frequency:** Since the update step happens less frequently, you might want to adjust the logging frequency to ensure you're still logging an appropriate amount of information.
- **Epoch Duration:** Note that while the effective batch size is larger, the number of updates per epoch will be reduced, which could affect the convergence behavior depending on your specific scenario.

## Example Integration

Here's how the full integration might look like with the modified trainer setup for using gradient accumulation:

python

```
if __name__ == "__main__":
    import os
    import wandb
    from pytorch_lightning.callbacks import ModelCheckpoint

    # Initialize Weights & Biases
    wandb.init(project="image-classification")

    # Define the data module
    data_dir = 'path_to_your_dataset'
    data_module = ImageDataModule(data_dir)

    # Define the model
    model = ResNetClassifier(num_classes=10)

    # Define the logger
    wandb_logger = WandbLogger(project="image-classification")
```

```
# Define the trainer with gradient accumulation
trainer = pl.Trainer(
    max_epochs=10,
    gpus=1 if torch.cuda.is_available() else 0,
    accumulate_grad_batches=4, # Here is the gradient accumulation
    logger=wandb_logger,
    callbacks=[ModelCheckpoint(monitor='val_loss')]
)

# Train the model
trainer.fit(model, data_module)

# Finish the Weights & Biases run
wandb.finish()
```

This setup should help you utilize your limited GPU more effectively by allowing for larger batch sizes through gradient accumulation.

Was this response better or worse?

ChatGPT can make mistakes. Check important info.